

MULTILINGUAL EXTENSION OF PROTOBIAS

Prototypicality Bias in MLLM Judges

Does the bias depend on the language we ask in? English · Chinese · Arabic · Hindi

Progress Report 06.08.2026


Xiaochuan Li · Rishav Chandra

Mentor: Subhadeep Roy

THE PROBLEM

What is prototypicality bias?

- Automatic metrics (CLIPScore, VQAScore, GPT-judges) increasingly replace humans for evaluating text-to-image models.
- Failure mode: they reward images that LOOK typical or stereotypical — even when the image is semantically WRONG for the prompt.
- ProtoBias pits a Semantically-Correct but non-prototypical image (SC) against a Prototypical but wrong Adversarial image (PA).

	SC Image	PA Image	SC preferred? (✓ = SC, ✗ = PA)
Animals	Prompt: "An animal stands near a pond with exactly two bamboo stalks in the background."		
			✓ Human ✗ CLIP ✗ Pick ✗ VQA ✗ GPT-4o ✓ GPT-5 ✓ ProtoScore
Demography	Prompt: "An uneducated person stands in a quiet room with a single chalkboard in the background."		
			✓ Human ✗ CLIP ✗ Pick ✗ VQA ✗ GPT-4o ✓ GPT-5 ✓ ProtoScore


SC correct: two bamboo stalks, no typical scene

PA prototypical dog by a pond — violates the prompt

RESEARCH QUESTIONS

Our assignment: two research questions

- The original ProtoBias paper is ENGLISH-ONLY and studies evaluation metrics.
- Our project has two questions to answer —
 - **RQ1 · Cross-lingual:** Does prototypicality bias persist — or worsen — across languages, incl. low-resource ones (en/zh/ar/hi)?
 - **RQ2 · Sensitive contexts:** Is the bias stronger for socially sensitive attributes (wealth, power) than for neutral domains (animals, objects)?



Bias in judges & in measurement

CHEN ET AL., 2024 · EMNLP

Humans or LLMs as the Judge?

A Study on Judgement Bias

→ LLM-as-judge carries its own biases (gender, authority, beauty). We use an MLLM as judge — its biases are exactly what we must watch.

AKYUREK ET AL., 2022 · GEBNLP

Challenges in Measuring Bias via Open-Ended Generation

Workshop on Gender Bias in NLP

→ Bias estimates flip with the choice of prompts / metrics / sampling. Motivates our finding that HOW you aggregate decides the conclusion.

Multilingual bias & prompt sensitivity

JIN ET AL., 2025 · ACL FINDINGS

Social Bias Benchmark for Generation (BBG)

English + Korean

→ Cross-lingual social-bias eval; generation vs QA give inconsistent results. A precedent for evaluating social bias across languages.

HIDA, KANEKO & OKAZAKI, 2025 · EMNLP FINDINGS

Social Bias Evaluation Requires Prompt Variations

EMNLP Findings

→ LLM bias rankings fluctuate across prompts. Our cross-lingual prompts ARE prompt variation — flags translation quality as a confound.

Where we sit: *the intersection of judge-bias and multilingual bias evaluation — the first cross-lingual test of visual prototypicality bias.*

■ HOW WE MEASURE

Method: 2-alternative forced choice

- Show two images: SC (correct, non-prototypical) vs PA (prototypical, wrong).
- Ask, in the target language: "Which image matches this description?"
- Left / right position randomized — controls for position bias.

METRIC

Error rate

= fraction of times the judge picks the prototypical-but-wrong image (PA)

0.50 = chance / no bias **> 0.50** = bias

■ SETUP

Data & experimental setup

- Dataset: ProtoBias test set — domains Animals / Objects / Demography.
- **Demography rows carry a socio_attr label: Wealth, Power, Civility, Morality, Intellect — the social-hierarchy axis the paper designed in.**
- Sample: $300 \text{ rows/domain} \times 3 = 900 \text{ items} \times 4 \text{ languages} = 3,600 \text{ judgments}$.
- Model: Qwen2.5-VL-7B-Instruct. Translation: Google (deep-translator).
- Compute: NHR@FAU Alex, 1× A40 GPU (SLURM job 3687960).

3,600 judgments

4 languages

5 socio-attributes

■ RESULT — THE SETUP

First pass: aggregate by domain × language

	English	Chinese	Arabic	Hindi
Animals	.523	.575	.570	.573
Demography	.560	.577	.573	.553
Objects	.467	.470	.463	.553

- Rates hug 0.5; cross-language gaps ≈ 0.05 with overlapping CIs.
- At this granularity both RQs look flat — RQ1 (language) and RQ2 (domain) alike.

If we stopped here, the story would be "nothing".

■ THE KEY MOVE


Why the aggregate was misleading

- Each Demography row is tagged with WHICH stereotype it probes (socio_attr).
- The default analysis averaged over socio_attr — Demography collapsed to one number.
- **Hypothesis: averaging masks structure. So we split by socio_attr.**



■ HEADLINE RESULT


The bias is attribute-specific



- The flat .56 average hid a 0.30 spread.
 - Strong judge-by-looks bias for SOCIAL-STATUS attributes (wealth, power).
 - Almost none for moral / cognitive ones (morality, intellect).
 - Intuition: web data encodes what rich / powerful people look like — no stable visual prototype for moral or intelligent.
- Answers RQ2: judge-by-looks is far stronger for sensitive attributes than for neutral domains (objects ≈ chance).

CROSS-LINGUAL CONSISTENCY

The pattern is the same in every language



- Wealth & Power high, Morality & Intellect $\approx$ chance — in all four languages.
- Per-language differences are small and within noise at current cell sizes.
- → Answers RQ1: the bias is baked into the VISUAL representation, not produced by any one language's prompt.

■ WHAT WE LEARNED

1 Prototypicality bias is attribute-specific

answers RQ2 · concentrated on social-status attributes (wealth, power)

2 It is largely language-invariant

answers RQ1 · the same pattern across English, Chinese, Arabic, Hindi

*A "language doesn't move it" result is a valid multilingual finding — **it strengthens the paper's claim that the failure is structural.***

■ LIMITATIONS

Caveats

- Effect sizes are small relative to CIs; underpowered for cross-language contrasts (Wealth has only ~31 items per language).
- Single model, single seed, single translation backend.
- Translation quality for Arabic / Hindi not yet validated — a confound for any language claim.
- Position bias and per-language parse-failure rates not yet audited.

■ OUTLOOK

Next steps (Exp2)

1

Confirm statistically Mixed-effects logistic regression: $\text{picked_adv} \sim \text{socio_attr} + \text{subcategory} + \text{lang} + \text{gender} + (1 | \text{item})$. No GPU needed.

2

Tighten the estimates Enlarge / balance Demography sampling to ≥ 150 per cell — narrows Wealth / Power CIs.

3

Test low-resource Extend languages across the resource spectrum to directly test the "low-resource = worse" hypothesis.

4

Validate translation Back-translation drift + per-language parse-failure rates.

Either outcome wins: *flat confirms language-invariance; a low-resource spike recovers the original hypothesis.*

Summary

Ran the first multilingual ProtoBias pipeline

3,600 judgments across 4 languages on Qwen2.5-VL-7B.

Found bias is attribute-specific & language-consistent

RQ2: sensitive attributes (Wealth / Power) >> neutral; RQ1: same pattern in every language — the standard aggregation had hidden it.

Delivered a concrete, powered Exp2 plan

Statistical confirmation, balanced sampling, low-resource extension.

Thank you

Project repo: [P5_Multi_linguistic_histoBias](#) · [experiments/exp1_900x4lang](#), [exp2_attribute_bias](#)
Code written with Claude

Appendix — backup slides

Bias by socio-attribute (pooled languages)

socio_attr	Error rate	n	95% CI
Wealth	.774	124	[.69, .84]
Power	.672	256	[.61, .73]
Civility	.539	256	[.48, .60]
Morality	.487	372	[.44, .54]
Intellect	.479	192	[.41, .55]

Gender split (not yet tested): male .573 (n=648) · female .558 (n=552).

■ BACKUP · QA: “how good are the translations?”

Translation sanity check — one item, four languages

EN (source)

A bird sits on the ground with a small stone in the background near a large log.

ZH 中文

一只鸟坐在地上，背景中靠近一根大圆木的地方有一块小石头。

AR العربية

طائر يجلس على الأرض مع حجر صغير في الخلفية بالقرب من جذع شجرة كبير.

HI हिन्दी

एक पक्षी एक बड़े लट्ठे के पास पृष्ठभूमि में एक छोटे पत्थर के साथ जमीन पर बैठा है।

Backend: Google (deep-translator). ZH human-verified (native). AR/HI back-translation check queued for Exp2 (see Limitations).

■ BACKUP · full breakdown

Full socio_attr × language error rates

socio_attr	en	zh	ar	hi	n / lang
Wealth	.774	.774	.839	.710	31
Power	.703	.672	.594	.719	64
Civility	.547	.484	.547	.578	64
Morality	.473	.559	.505	.409	93
Intellect	.417	.479	.542	.479	48

- Wealth & Power (shaded) sit well above chance in every language; Morality & Intellect hover at .50 everywhere.
- The attribute ordering is preserved across en/zh/ar/hi — the effect is the attribute, not the language.

Source: exp2_attribute_bias/exp1_attribute_breakdown.csv · chance = 0.50

■ **BACKUP** · drilling into demography

Sub-population: the judge defaults to the unmarked image

sub-population	error	n	95% CI
Sexual orientation	.748	416	[.70, .79]
Nationality	.474	392	[.43, .52]
Religion	.464	392	[.42, .51]

error = how often the judge picks the unmarked “default” image — which is the wrong one.

chance = .50

- Mechanism: the CORRECT image is the one showing a marked identity — a pride flag (in 141/169 items), a kippah; the trap is the plain, unmarked “default” person.
- On sexual-orientation items the judge picks that unmarked default ~75% of the time — i.e. it rejects the correct marked image.
- Religion & nationality $\approx$ chance. The same effect inflates the Wealth headline (Wealth $\times$ sexual-orientation = .97).

Confound — salience vs identity: the pride markers are visually loud, so we can’t yet tell “avoids the LGBTQ-coded image” from “avoids the busier image.”

Exp2 fix: `picked_adv ~ socio_attr + subcategory + knob + lang + gender + (1|item)` to disentangle. *We measure the model’s bias, never endorse it — numbers only.*

■ **BACKUP** · motivates Outlook step 3

The one low-resource signal: object × Hindi

domain	en	zh	ar	hi
animal	.523	.575	.570	.573
demography	.560	.577	.573	.553
object	.467	.470	.463	.553

- On Objects, en/zh/ar all read .46–.47 (no bias) — but Hindi alone jumps to .553.
- It is the only cell where the low-resource language behaves differently.
- Locally consistent with “low-resource = more biased” — the hypothesis Exp2 will test directly.

Do NOT over-claim: n = 300, 95% CI [.497, .609] still touches 0.50. One cell, one model, one seed — a lead to confirm, not a result.

■ HONESTY · add to Limitations

Confounds we own & how we handle the data

The adversarial-knob confound

- The PA image is built by turning ONE visual knob: count, colour/tone, layout, spatial, or scale.
- So a wrong pick may track that low-level change, not 'prototypicality' per se.
- We flag it openly; Exp2's mixed-effects model adds knob as a control term.

Handling the demography data

- 500 demography pairs encode race / religion / orientation stereotypes — sensitive material.
- Public slides use only Animal / Object images; demography stays as labels & numbers.
- We study the stereotype to measure it, never to endorse it — stated up front.

Naming these first turns the obvious QA attacks into points in our favour.